

Publications

(making it seems more impressive than it is because it's all I got)

- [1] **Hieu M Vu** and Tan Minh Nguyen. “Angular Steering: Behavior Control via Rotation in Activation Space”. In: *Advances in Neural Information Processing Systems* (2025). **Spotlight (top 3.5%)**.
- [2] Ryan Lee*, Dung Viet Nguyen*, **Hieu M Vu***, Lei Zhang, Linh Duy Tran, and Tan Minh Nguyen. *Momentum Steering: Activation Steering Meets Optimization*. Under review at ICLR 2026. 2025.
- [3] Dung V Nguyen*, **Hieu M Vu***, Nhi Y Pham*, Lei Zhang, and Tan M Nguyen. *Activation Steering with a Feedback Controller*. Under review at ICLR 2026. 2025. arXiv: 2510.04309 [cs.LG].
- [4] Nguyen Hong Son, **Hieu M Vu**, Tuan-Anh D Nguyen, and Minh-Tien Nguyen. “Jointly Learning Span Extraction and Sequence Labeling for Information Extraction from Business Documents”. In: *2022 International Joint Conference on Neural Networks (IJCNN) Oral*. IEEE. 2022, pp. 1–8.
- [5] Tuan-Anh D Nguyen, **Hieu M Vu**, Nguyen Hong Son, and Minh-Tien Nguyen. “A Span Extraction Approach for Information Extraction on Visually-Rich Documents”. In: *Document Analysis and Recognition-ICDAR 2021 Workshops: Lausanne, Switzerland, September 5–10, 2021, Proceedings, Part II 16. Best Paper Award*. Springer. 2021, pp. 353–363.

Education

(grade is just a number and I got lucky)

Bachelor Degree, Computer Science (Honours Programme)

2020

UET - Vietnam National University, Hanoi

GPA: 3.83/4.00 (Rank: 1st/600+)

- Highest Ranking Graduate.

Experience

(trying to up-sell myself)

Visiting Scholar

Sept 2025 — Oct 2025

National University of Singapore, Advisor: Prof. Tan Nguyen

Fully funded visit to NUS to collaborate on research with Prof. Tan Nguyen.

AI Research Engineer

July 2024 — Present

Torilab

- Building AI avatars.

- Conducting research on controlling LLMs' behaviours without prompting (i.e. activation steering).
- LLMs fine-tuning (SFT, DPO, classification), prompt engineering, LLMs serving (vLLM, sglang, LitServe) with dynamic batching.
- Developing the core framework for the customizable chatbots with memory, tool use, dynamic prompts and multi-turn interactions.
- Other random tasks: User's engagement detection with head-pose tracking; Image Generation with Diffusion-based models and LoRAs; multi-speaker diartization.

AI Research Engineer

Nov 2018 — May 2024

Cinnamon AI

- Developed RAG-based applications.

- Co-creator of **kotaemon** (24.5k+ ★): An open-source tool for local RAG application built for both end users and developers.
- Built demos for LLM-powered applications targeting the Insurance domain.

- Researched, developed, and implemented AI solutions for Document Image Understanding.

- Information Extraction and Cross-lingual adaptive pre-training for a low-resource language (Japanese).
- Created new technical assets by introducing new Information Extraction models that became the new standard for client projects. Increased the f1-score by 2% - 7%.
- Publish **papers** on Information Extraction at peer-reviewed conferences, 1 Best Paper Award at DIL-ICDAR'21.

- Developed data-driven products and processes.

- Implement synchronization and local version control for the internal data management system.
- Initiated data standardization, defined and implemented processes regarding data life cycle and organization.
- Other tasks: Support cross-department alignment; Conduct/Facilitate technical sharing sessions; Training/Mentoring.

Undergraduate Research Assistant

Aug 2017 — Sep 2019

IOT Lab, UET - Vietnam National University

- **Predictive models for wellbore data using machine learning and statistical methods.**
 - Facies/rock type classification; Time-series Analysis; Permeability Regression; Integrated Prediction Error Filter Analysis (INPEFA) curve calculation; Cumulative and Federated Learning.

Honours and Awards

(more uncomfortably showing off)

Best Paper Award

Sep 2021

ICDAR 2021, Workshop on Document Images and Language

Paper title: [A Span Extraction Approach for Information Extraction on Visually-Rich Documents](#)

Accepted for oral presentation and awarded the Best Paper Award at Workshop on Document Images and Language, ICDAR 2021.

Certificate of Highest Ranking Graduate

Aug 2020

UET - Vietnam National University

Awarded to students who graduated with the highest GPA amongst the graduating class.

Top 4 Zalo AI Challenge 2018 - Voice Track (Individual participant)

Sep 2018

Zalo, VNG Corporation

Finished at 4th place on the Private Leaderboard of the Voice Gender/Accent Classification challenge.

[Zalo AI Challenge](#) is an annual Kaggle-like competition hosted by Zalo - one of the biggest tech companies in Vietnam. In 2018, the competition attracted over 700 teams competed in 3 challenges.

Personal Projects

(probably no one cares but can't waste empty space)

Kotaemon (24.5k+ ★) - An open-source tool for local RAG application.

Jan 2024 — May 2024

Open-source project, Co-creator



- A local RAG-based tool for chatting with your documents. Built with both end users and developers in mind.
- For end users: A local Question Answering UI for RAG-based QA.
- For developers: A framework for building your own RAG-based QA pipeline.

Data Utility Improvement Experiment for DECAF

Oct 2022 — Nov 2022

Personal research



- A personal research on Causal Inference, Algorithmic Fairness and specifically the paper [DECAF: Generating Fair Synthetic Data Using Causally-Aware Generative networks](#).
- Improved data utility of the DECAF method using alternating graph during synthesis while still achieving similar level of fairness.

Gender/Accent Classification for Vietnamese short voice recordings

Aug 2018 — Sep 2018

Zalo AI Challenge 2018



- Classify the speaker's voice in a recording (typically under 3 seconds) by gender and regional accent.
- **4th place** on the Private Leaderboard, achieved **79.208% accuracy** in 10 days as an **individual participant**.

Skills

(for those bots that love keywords)

Techincal Fields	LLMs, ML Interpretability, Information Extraction, Document Understanding, NLP, Computer Vision, Image Processing, Time-series Analysis.
ML/AI Development	Pytorch , Transformers , RAG , LangChain, Haystack, LlamaIndex, Tensorflow/Keras, Scikit-learn, OpenCV.
Software Development	Git , Github Action , Docker, CircleCI, DVC.
Programming Languages	Python , C/C++, Java, Shell Script.
Industrial Domains	Insurance, Manufacturing, Virtual Companion.
Environments	GCP, AWS, Linux, Windows/WSL.
Natural Languages	Vietnamese (native), English (IELTS 7.5), Japanese (JLPT N4).
Other	Research (Google Scholar), Problem Solving (Leetcode).